

AI Engineer Intern Assignment 1

Name: Parul Sharma

Topic: LLM Comparison

Submitted to: EDXSO

Date: 13/01/2026

Rank 1 – GPT-4 / GPT-4o (OpenAI)

Advantages

1. It explains things step by step, When we're stuck, that kind of guidance makes a huge difference.
2. For coding—whether we're writing scripts, debugging, or just trying to understand what's broken—it's honestly rock solid.
3. It keeps track of the conversation, even if we're bouncing around or breaking things up. We don't have to repeat ourself.

Disadvantages

1. The best features sit behind a paywall.
2. It's closed source, so we can't tinker with the model itself.
3. Sometimes it just jumps right into advanced stuff and leaves we in the dust if we're not ready.

Comfort Level: Advanced

Justification: I put GPT-4/GPT-4o at the top because I trust it most for serious academic work or building real projects. The reasoning and accuracy are just hard to beat.

Rank 2 – Claude (Anthropic)

Advantages

1. It's great at taking long, dense documents and boiling them down to what matters.
2. The answers are organized and easy to follow.
3. With sensitive or difficult topics, it handles things carefully instead of just blurting out whatever.

Disadvantages

1. It's not as strong on technical coding as GPT-4.

2. If we want wild or really creative answers, Claude's not the one.
3. It doesn't play as nicely with outside tools compared to some others.

Comfort Level: Intermediate

Justification: I put Claude second because it's great at reading, summarizing, and explaining. But for deep technical help, it can't quite keep up with GPT-4.

Rank 3 – Gemini (Google)

Advantages

1. If we're looking for reliable facts or general info, Gemini does the job.
2. It fits right into Google's world. Simple, fast, reliable.
3. Can understand text and images at the same time, which is super practical.

Disadvantages

1. Sometimes it just gets it wrong, or misunderstands what we're trying to say — it really depends on how we ask and what we ask.
2. Some answers feel a bit bland, like it's replying just to reply, without giving anything deep or truly helpful.
3. If we need deep, careful reasoning, this isn't our best choice.

Comfort Level: Intermediate

Justification: I put Gemini in third place. It's really good for daily use, but when we need deep answers or more reliable accuracy, it can feel a bit inconsistent.

Rank 4 – LLaMA (Meta)

Advantages

1. Since it's open-source, it gives full freedom to explore it, change things, and customize it the way needed.
2. It's really good for research work, and also fun to use when just experimenting with new ideas and testing different things.
3. The best part is that it can be run directly on a personal PC/laptop, which is great because it gives complete control and doesn't force dependence on cloud platforms.

Disadvantages

1. Setting it up properly needs some technical knowledge, otherwise it can be confusing and messy.
2. Right after installing, it may feel less powerful compared to big commercial models, and usually needs tuning to give better results.
3. For beginners, it can feel confusing at first because there's a learning curve and a lot of new things to learn.

Comfort Level: Beginner to Intermediate

Justification: I ranked LLaMA fourth. It's amazing for researchers and people who love experimenting, but some setup and tweaking is needed to get the best performance out of it.

Rank 5 – Mistral

Advantages

1. It's light and fast—doesn't need a lot of resources.
2. Perfect for devices or situations where we don't have much power to spare.
3. The open models let developers do pretty much whatever they want.

Disadvantages

1. It's not built for heavy reasoning or complex problems.
2. The community is smaller, so there's less help if we get stuck.
3. Not really made for big, diverse projects.

Comfort Level: Beginner

Justification: I put Mistral last because it works best for specific, performance-focused jobs, not as a general-purpose AI we had use for everything.